

A. Obtain representations; fix their component groups




D_1, D_2 : visual features; Q_1, Q_2 : caption features; V : image latents.

B. Measure which components help a conditional task

Image + “The [MASK] gives the cup.” → predict the missing word.


D_1 mostly noise

D_1 clearer




Hold other inputs and target noise fixed; mask the caption before encoding.

C. Dependencies guide timing



Curves illustrate overlapping schedules.

D. Dependencies guide module access



Turn the head off; restore its output.
Does the benefit of D_1 disappear
and return?

Repeat across noise levels and tasks.

Jointly learn schedules and transformer weights; representation groups remain fixed.